

Recommender Systems Project


Movie Recommendations with Graph Neural Network

Group:

22111013 – Atul Kumar

22111054 – Shubham Rathore

22111069 – Vivek Kumar Gautam



Problem Statement


Yield recommendations on the MovieLens-100K dataset

About the dataset

- 100,000 ratings (1-5) from 983 users on 1682 movies
- Each user has rated at least 20 movies
- Some demographic information is present about the users
- Collected on the MovieLens website during Sept '97 through Apr '98
- **Citation:** F. Maxwell Harper and Joseph A. Konstan. 2015. The MovieLens Datasets: History and Context. ACM Transactions on Interactive Intelligent Systems (TiiS) 5, 4, Article 19 (December 2015), 19 pages.
DOI=<http://dx.doi.org/10.1145/2827872>


Recommender Systems as a Graph

- Any recommendation system can be naturally modeled as a **bipartite graph**
- Nodes:
 - Users
 - Movies
- Edges:
 - An edge denotes an **interaction** between a user and an item
 - User interacted with a movie by giving it a good (or a bad?) rating



Recommendation Task

- **Given:** Ratings from users on movies
- Predict which movie the user will like in future **i.e., predict an edge**
- We need to get a score $f(u, v)$, for $u \in U, v \in V$
- This score will indicate the probability of a user liking a movie




Top-K Recommendation

- For each user, we recommend K items
- For recommendation to be effective, K needs to be much smaller than the total number of items
- K is typically in the order of 10 – 100
- Include as many **positive items** as possible in the top- K recommended items
- Positive items = Items that the user will interact with in the future

Evaluation Metric: Recall@K

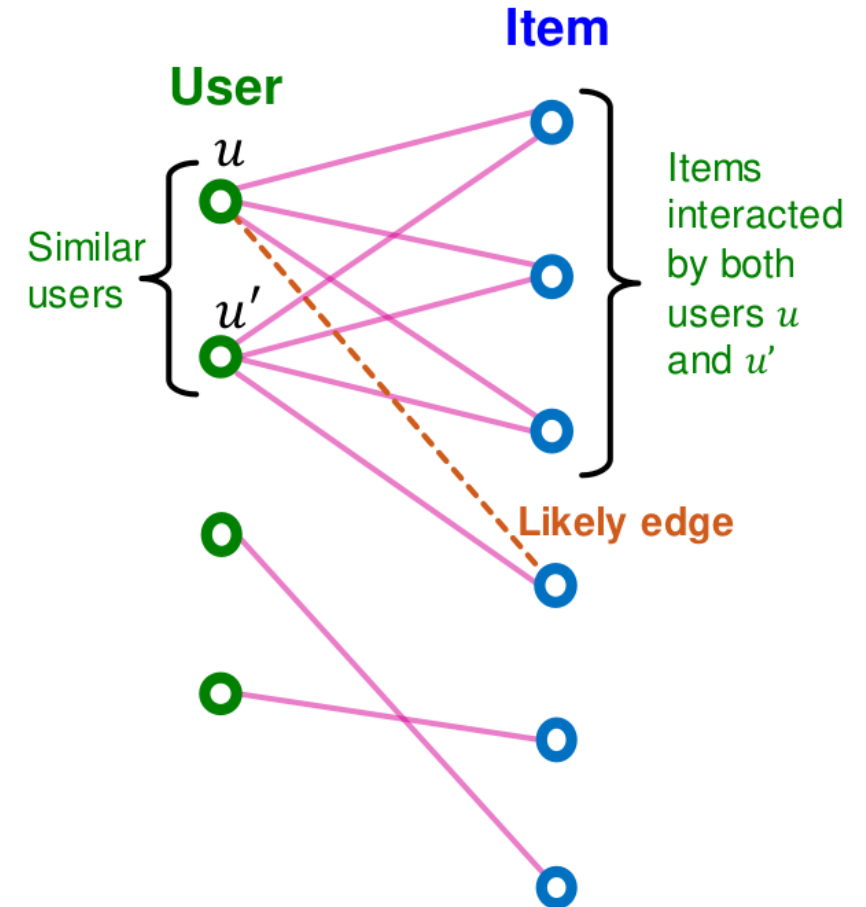
- For each user u ,
 - Let P_u = set of positive items the user will interact with in the future,
 - Let R_u = set of items recommended by the model
 - $|R_u| = K$, since the model recommends only the top-K items
 - Recall@ K for user u is $(P_u \cap R_u) / P_u$



$$\text{Recall@K} = \frac{1}{|\text{Users}|} \sum_{u \in \text{Users}} \left(\frac{P_u \cap R_u}{P_u} \right)$$

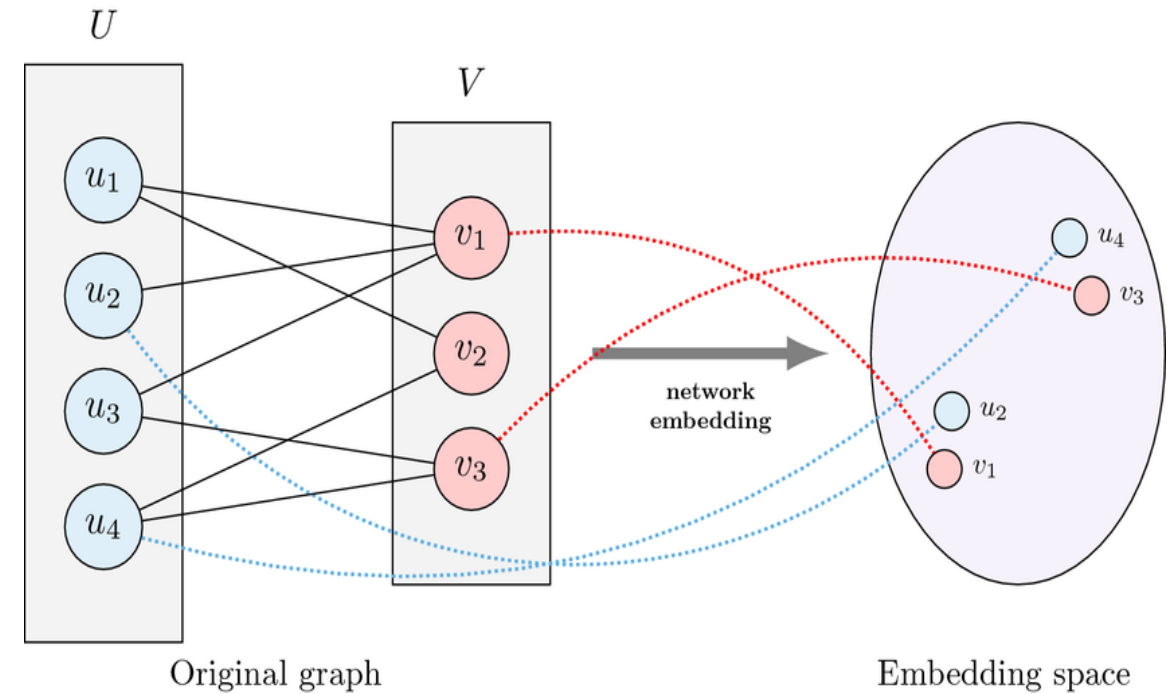
Collaborative Filtering

- Recommend movies to a user based on the preferences of similar users
- Similar users tend to prefer similar movies – collaborative filtering idea
- **Task:** Capture similarity between users/movies



Which model to use?

- A complex feature-based model could simply memorize the user-item interactions without learning similarity
- Solution: **Embedding-based models**
- Low-dimensional embeddings are forced to capture similarities to fit the data
- Better predictions on **unseen** user-item interactions



Training Objective

- Optimize the model to **achieve high recall@ K on seen items** i.e., the training items
- We hope that this will lead to high recall@ K on unseen (test) items

Problem with recall@ K

- It is not differentiable – so we **cannot apply a gradient-based optimization**
- However, another loss function called the Bayesian Personalized Ranking (BPR) loss can be used
- It aligns with the recall metric and is differentiable

So Far

- Use Recall@K as a metric for personalized recommendation
- Use an embedding-based model and learn -
 - User encoder (to generate user embeddings)
 - Item encoder (to generate item embeddings)
 - Score function (to predict user-item interaction likelihood)
- BPR loss for achieving the optimization objective

Choosing a GNN architecture

- GNNs* have been widely successful in recommendation tasks
- **NGCF** [Wang et al. 2019]
 - Neural Graph Collaborative Filtering (NGCF)
 - Incorporates non-linear activations and complex feature transformations
 - Learns graph-aware user-item embeddings
- **LightGCN** [He et al. 2020]
 - Light Graph Convolutional Network (LightGCN)
 - Empirically proves that NGCF's complex methods are not required for good performance
 - Explicitly models the graph structures
 - Assumes no user/item features

*Sanchez-Lengeling, et al., "A Gentle Introduction to Graph Neural Networks", Distill, 2021.

LightGCN Overview

- LightGCN has the following convolution/propagation rule between each layer:

Embedding of user u_m at layer $k+1$

$$\boxed{u_m^{(k+1)}} = \sum_{i_n \in N(u_m)} \frac{1}{\sqrt{d_{u_m}} \cdot \sqrt{d_{i_n}}} \cdot \boxed{i_n^{(k)}}$$

Embedding of item i_n at layer $k+1$

$$\boxed{i_n^{(k+1)}} = \sum_{u_m \in N(i_n)} \frac{1}{\sqrt{d_{i_n}} \cdot \sqrt{d_{u_m}}} \cdot \boxed{u_m^{(k)}}$$


Degree of node u_m

Degree of node i_n


- The only trainable model parameters in LightGCN are the 0-th layer embeddings of all users and all items

Training and Evaluation


- A 3-layer LightGCN model takes ~3 minutes to train for 50 epochs (on Nvidia T4 GPU)



Training and Evaluation



*Recall@10 and Precision@10
~0.22 and ~0.31 respectively*



*Recall@20 and Precision@20
~0.35 and ~0.24 respectively*

Generating Recommendations

➞ some movies that the user rated highly:

Toy Story (1995)
Twelve Monkeys (1995)
Seven (Se7en) (1995)
Apollo 13 (1995)
Star Wars (1977)
Fugitive, The (1993)
Terminator 2: Judgment Day (1991)
Silence of the Lambs, The (1991)
Rock, The (1996)
Die Hard (1988)

top 10 recommended movies for the user:

Father of the Bride (1950)
Supercop (1992)
Full Metal Jacket (1987)
Things to Do in Denver when You're Dead (1995)
Clockwork Orange, A (1971)
Crow, The (1994)
Pink Floyd - The Wall (1982)
Strange Days (1995)
Bonnie and Clyde (1967)
Wings of Desire (1987)



THANK YOU